

Ankit Raj

Forward Deployed Engineer - Full-Stack, Edge & AI-Native Builder

Stockholm, Sweden | rajankit749@gmail.com | ankitraj.cloud | linkedin.com/in/raj-ankit | github.com/restlessankykyy

PROFESSIONAL SUMMARY

Builder-at-heart engineer with 10+ years shipping production-grade systems, not just advising on them. I live in the IDE with a ship-first mentality - decomposing ambiguous, zero-to-one problems into scalable, resilient architectures and owning solutions end to end. Full-stack depth from frontend performance and backend logic to distributed infrastructure, serverless/edge, networking, and security primitives across AWS, Azure, and GCP. AI-native in my daily workflow: I build with LLMs, agents, and MCP tooling and ship open-source developer tools used by the community. I've owned mission-critical services with real users - on-call, incident response, and the rigor of five-nines uptime (99.9%). High agency, comfortable embedding directly in partner engineering teams and engaging VP+ stakeholders, translating complex engineering trade-offs into business impact while staying deeply technical.

CORE COMPETENCIES

Production Engineering	Production-Grade Code, Mission-Critical Services, On-Call & Incident Response, Five-Nines Uptime, High-Velocity Shipping
Full-Stack	TypeScript, Python, Node.js, Go, JavaScript - Frontend Performance, Backend/APIs, Database Optimization, Distributed Systems
Edge & Serverless	Serverless (Lambda, Azure/Cloud Functions), Edge Patterns, API Gateways, Cloud Run, Event-Driven Architecture, CDN/DNS Concepts
Networking & Security	Zero Trust, VPC/Networking Primitives, IAM, OAuth 2.0 / OIDC, Rate Limiting, Observability (ELK, Grafana, Prometheus)
AI-Native Development	AI-Augmented Workflows, LLM/GenAI, Agentic AI & MCP Tooling, RAG, Rapid Prototyping, Copilot/Claude-Driven Delivery
Cloud Ecosystem	AWS, Microsoft Azure, Google Cloud - Serverless, Networking, Security, Kubernetes (AKS/GKE), Docker
Systems Design	Decomposing Complex Requirements, Scalable & Resilient Architecture, Big-Picture + Implementation Detail, Infrastructure as Code (Terraform)
Ways of Working	High Agency, Zero-to-One, Embedded Partner Delivery, Standups & Design Reviews, Executive Presence (VP+), Open-Source & Blogging

PROFESSIONAL EXPERIENCE

Lead Engineer & Architect - AI & Multi-Cloud 2024 - Present
H&M Group, Stockholm

- Build production AI systems hands-on - multi-agent workflows with MCP tooling and RAG for incident response and security governance on Microsoft Foundry and GCP Vertex AI, shipping real capabilities for 200+ product teams
- Embedded as the technical lead across engineering teams, decomposing ambiguous requirements into scalable, resilient architectures and owning solutions end to end with a ship-first mindset
- Engineer serverless and edge-style integrations (Functions, API gateways, event-driven flows) across AWS, Azure, and GCP, going deep on networking, security, and observability primitives
- AI-native in daily delivery - integrated LLM- and agent-augmented workflows to accelerate prototyping and shipping; evaluated OpenAI, Claude, and Gemini for developer experience
- Translate complex engineering trade-offs into business impact for VP+ stakeholders, driving a 35% cloud cost reduction and high adoption signals across the estate

Lead Software Engineer - Platform & Developer Tooling 2023 - 2024
H&M Group, Stockholm

- Designed and shipped an Internal Developer Platform with self-service provisioning - EventGrid + Azure Functions event-driven architecture (TypeScript) reducing onboarding from days to minutes
- Owned mission-critical platform services on Azure & Google Cloud with 99.9% uptime, participating in incident response and operational rituals for 200+ teams
- Stayed an active practitioner in the IDE - evaluated and integrated OpenAI, Claude, and Gemini, shipping developer-productivity tooling that improved team velocity by 40%
- Ran design reviews and surfaced reusable patterns, lifting standards across teams while maintaining high shipping velocity

Forward / Embedded Engineer - FSI Modernisation 2019 - 2020
Cognizant - Ikano Bank, Stockholm

- Embedded directly within a financial-services client's engineering team, writing production code to migrate AML & KYC APIs from an on-prem gateway to a serverless, event-driven platform - 70% latency reduction
- Built and deployed production infrastructure (Lambda, API Gateway, Cognito, DynamoDB, Step Functions) handling 500K+ daily transactions, with Terraform IaC and CloudWatch observability
- Shipped an LLM-powered autonomous migration agent for automated code and policy conversion - cutting delivery time by 60% and demonstrating AI-native, ship-first delivery
- Owned the account's technical success end to end, identifying problems proactively and driving alignment with senior stakeholders under shifting priorities

Senior Software Engineer - Distributed Systems 2020 - 2023
H&M Group, Stockholm

- Built high-throughput GraphQL & REST APIs (50K+ RPM) and 15+ microservices with sub-100ms p99 latency on a Zero Trust, multi-cloud architecture (Azure gateways on GKE)
- Owned production services with real users - on-call, incident response, and architectural rigor to sustain reliability at scale
- Led CI/CD standardisation with GitHub Actions across 200+ repositories - 95% pipeline success rate, raising engineering standards across teams
- High agency: identified emerging reliability risks early through root-cause analysis and mobilised teams to act

DevOps / Platform Engineer Consultant - Automotive 2019 - 2020
Volvo Cars, Gothenburg

- Embedded with engineering teams to build and deploy an observability platform (ELK, Grafana, Prometheus) on AKS monitoring 100+ microservices, with real-time incident response
- Built CI/CD pipelines for OpenShift & Kubernetes - reduced deployment cycle time from hours to 15 minutes
- Authored runbooks and incident procedures - achieving 99.5% SLA compliance in production

Cognizant Technology Solutions

- Built production cloud applications and CI/CD pipelines (CodePipeline, Docker, ECS/Fargate) for 5+ enterprise clients, operating directly within partner codebases
- Developed a React-based Cloud Management Dashboard with real-time infrastructure and cost analytics - full-stack delivery front to back
- Won 1st place at Enterprise AWS Game Day 2018 (200+ participants), demonstrating deep technical skill under production pressure

OPEN SOURCE & BUILDS

Rubin - Local AI Coding Assistant (Open Source, GitHub)

Privacy-first local AI coding assistant on Ollama - streaming chat, 14-tool agent mode, MCP server integration, and inline completions. Full-stack TypeScript + VS Code API. Demonstrates AI-native building and developer-tooling depth.

OpenClaw Tesla - AI Vehicle Agent (Open Source, GitHub)

AI agent plugin for the Tesla Fleet API with 21 voice-controlled actions (lock, climate, charging, sentry), OAuth 2.0 secured, auto-wake and multi-vehicle. TypeScript - real production API integration and edge-style automation.

AML/KYC Serverless Migration (Ikano Bank, FSI)

Production serverless migration with an LLM-powered autonomous code-conversion agent. Lambda, API Gateway, DynamoDB, Terraform. 70% latency reduction, 60% faster delivery.

Internal Developer Platform - Event-Driven (H&M Group)

Self-service provisioning platform - EventGrid + Functions notification workflows in TypeScript, standardising golden-path infrastructure and cutting onboarding from days to minutes.

CERTIFICATIONS

- AWS Certified DevOps Engineer - Professional
- AWS Certified Developer - Associate
- Microsoft Certified Professional (MCP) - Exam 533: Azure Infrastructure
- Microsoft Certified: Azure Fundamentals (AZ-900)
- Microsoft OpenHack: AI-Powered Knowledge Mining · Containers · Security, Compliance & Identity · DevOps
- Google India Udacity Scholar

EDUCATION

Bachelor of Technology - Computer Science & Engineering

Galgotias University, Greater Noida, India

Specialisation in Cloud Computing & Virtualisation. Projects: IoT Heart Rate Monitor, Big Data Analytics.

PUBLICATIONS & THOUGHT LEADERSHIP

I Gave an AI Agent Root Access to a Sandboxed Linux Box - AI agent experimentation (Medium)

Google Just Compressed AI to 3 Bits - TurboQuant KV cache compression deep-dive (Medium)

What is Kubernetes? - Container orchestration deep-dive (FAUN / Medium)

Building Apps on Google Cloud Run - Serverless deployment patterns (Medium)